

# Learning Algorithms

NEUR 531-603 Introduction to Computational Neuroscience

Curtis Baker

**Primary Reading:** Information Theory, Inference, and Learning Algorithms, by David MacKay  
mainly chapter 39; also some chapters 38, 44

This book is freely available to download:


<http://www.inference.phy.cam.ac.uk/mackay/itila/book.html>

Individual chapters:

<http://www.inference.phy.cam.ac.uk/mackay/itprnn/ps/>

# Neural Networks - a brief history

1950s-60s: McCollough-Pitts neuron; "feature-detector" neurons in optic tectum, A17



1960s-70s: Rosenblatt Perceptron: architecture (single-layer)


novelty at the time: learning; distributed memory; neural inspiration

Minsky & Papert critique, difficulty with multi-layer networks

1980s-90s: revenge of the neural networkers: back-prop, connectionism, etc

concurrent influences: neural plasticity, NMDA receptors;

Donald Hebb; David Marr; Rumelhart, Hinton, Sejnowski



90s, 00s: rise of the machines: machine learning (neural or not)

probabilistic models, statistical learning theory

# Neural Networks

## **audiences**

cognitive science - connectionist models

neural network modeling - from loose metaphors, to specific models  
(e.g., development of ocular dominance stripes in visual cortex),  
to a theoretical endeavour in its own right

computer science: robotics / machine intelligence, computer vision

## **example applications of "machine learning"**

pattern recognition

handwriting, fingerprints, faces, license plates

decoding

"mind-reading" with fMRI

model parameter-fitting

system identification

finding patterns in data

cluster analysis

data mining

dimensionality reduction / data compression

# Types of learning algorithms

## 1. **supervised** - *data = inputs, targets (from a “teacher”)*

classification / recognition

output is discrete / categorical: e.g., Rosenblatt Perceptron

regression / function approximation

output is analog / functional: e.g., parameter-fitting

## 2. **unsupervised** - *data = set of multivariate values (without any “teacher”)*

density estimation - clustering, EM algorithm, mixture of Gaussians

efficient coding of natural sensory info

decorrelation, PCA, ICA

## 3. **reinforcement**

output is an “action”, optimized to maximize a “reward”

no access to examples of optimal or correct responses

instead, “agent” must *discover* them

*common concepts throughout:*

can often (optionally) be cast as “neural networks”

can often use as metaphorical neural models, or as data analysis tools

optimization algorithms - minimize an “error function” or “objective function”

# Classification / recognition


classify stimuli into categories

must be robust to variability in stimuli

dataset:


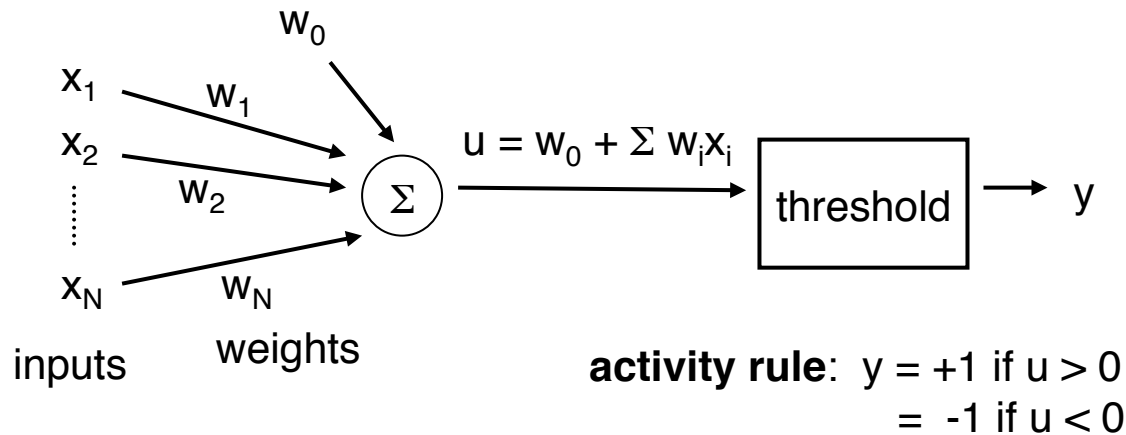
inputs ( $x$ 's), e.g. photoreceptors / pixels, or higher-level features / receptive fields  
corresponding classifications ( $T$ 's, i.e. the “teacher”)

**Rosenblatt Perceptron** - an early “binary classifier” -> classify into 2 categories



# Classification

**Architecture:** variables, and relationships between them



**learning (up-date) rule:** if correct ( $y=T$ ):

$$\Delta w_i = \eta x_i$$

$$\Delta w_0 = \eta$$

if incorrect ( $y \neq T$ ):

$$\Delta w_i = 0$$

*Feedback from teacher*

$T$  = "teacher"

$\eta$  ("eta") = learning rate parameter

**notes:** weights initially random  
sequential / on-line / stochastic mode  
sensitive to  $\eta$

**problems:** does not always converge; poor generalization; ad hoc (no underlying theory)

*Not generalizable*

( DEMO: Rosenblatt Perceptron )


# Classification with gradient descent

*principle:* specify an **error function**, which algorithm should try to minimize  
(in general, an “**objective function**”)

*dataset:*

inputs ( $x$ 's), e.g. photoreceptors / pixels, or higher-level features / receptive fields  
corresponding classifications ( $T$ 's, i.e. the “teacher”)

gradient descent:



$$\text{min: } \delta E / \delta w = 0 \quad \rightarrow \quad dw_i = -\eta \delta E / \delta w_i$$

*Slope*


$\eta$  (“eta”) = learning rate / “step size”

$E$  = error (objective) function  $\rightarrow$  gradient



# Classification with LMS gradient descent

**architecture**



**objective function:**  $E = 1/2 \sum (T_j - y_j)^2$  “least mean squares”


where  $y_j$  = network response, on trial  $j$

$T_j$  = “teacher”, i.e. desired or correct response

(for probabilistic model, with Gaussian noise, this objective function is optimal) ?

**activation function:** must be differentiable

-> logistic (sigmoid):  $y(u) = 1/(1 + e^{-u})$



-> **learning rule:**  $dw_i = -\eta \sum (T_j - y_j) x_i$

**notes:** for linear model, will always converge to unique minimum

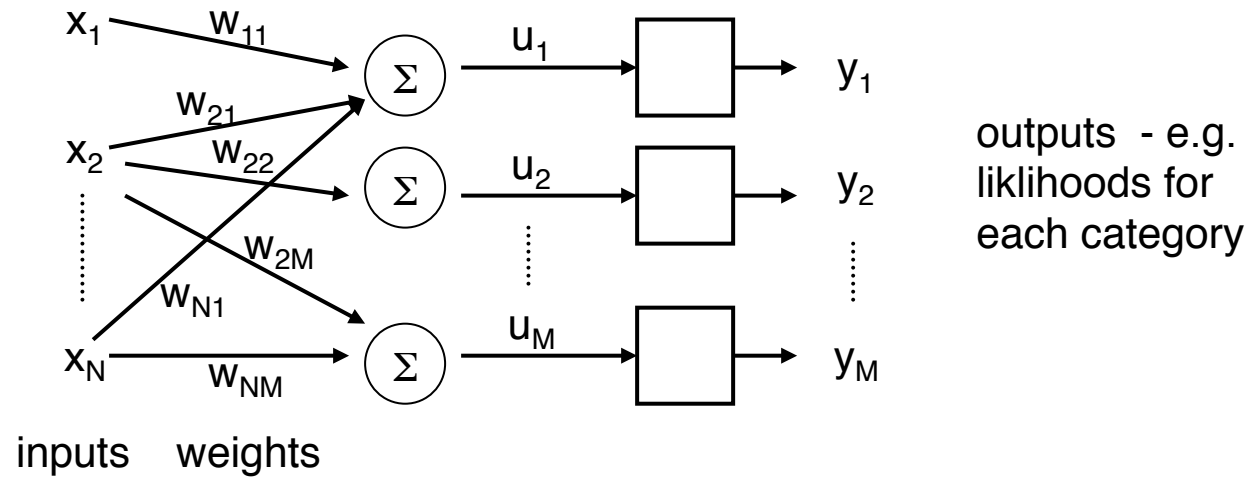
best run in batch mode

**problems:** only gives good classification if categories are linearly separable


( DEMO: binary classifier with gradient descent )

# Classification: extensions

more than two categories: **multi-output networks**



beyond linear separability: **multi-layer networks**



problems: multiple minima  
how many hidden units ?

**too many -> over-fitting**

# Regularization

**over-fitting** problem:

weights eventually diverge to very large values,  
giving only small improvements to error,  
while **degrading ability to generalize to new data**

-> see **MacKay, section 39.4 and Figure 39.5, 39.6** ( -> Assignment )

**regularization:** modify error function, to place a “penalty” on large weight values  
sometimes called “**weight decay**”

$$E = 1/2 \sum (T_j - y_j)^2 + \alpha \sum w_i^2$$

*For decay*

$\alpha$  = hyperparameter (not part of “activity rule” or model architecture -  
it is a parameter of the learning algorithm)

# Regression / function approximation

**general regression problem:**  $y = f(x)$

given input vector,  $x$ , and vector of outputs,  $y$ ,

-> find mapping function,  $f$

**compare to classification:**

outputs are analog, not categorical

“teacher”: try to optimize prediction of  $y$ -values

**applications:**

find best-fitting parameters of a model  
system identification

**linear regression:**  $y = w \cdot x$



dataset:  $x$  = inputs:  $x_1, x_2, \dots, x_N$   
 $y$  = outputs:  $y_1, y_2, \dots, y_M$

-> find  $w$  = weights:  $w_{11}, w_{12}, \dots, w_{1M}, \dots, w_{N1}, \dots, w_{NM}$

# Linear Regression

simplest case, single output:  $y = \sum w_i x_i$

**architecture**



**error function:**

$$E = \sum (y_i - \text{predicted } y_i)^2$$

gradient descent:

$$\text{min: } \delta E / \delta w = 0 \quad \rightarrow \quad dw_i = -\eta \delta E / \delta w_i$$

$\eta$  (“eta”) = learning rate parameter

**learning rule:**

$$dw_i = \eta (y_i - \text{predicted } y_i) x_i$$

a “*general linear model*” (GLM) – guaranteed to have a unique minimum

?

( DEMO: linear regression with gradient descent )

# Regression: extensions

**faster algorithms** to find optimum

- e.g. scaled conjugate gradient

**basis functions** (“features”) on front-end, e.g.:

- Gaussian weightings

- Gabor wavelets

**regularization**



# Regularization

**over-fitting** problem: weights eventually diverge to very large values, giving only small improvements to error, while degrading ability to generalize to new data

simple example: fitting a **high-order polynomial** to a *small* number of data points

-> curve fits those few points extremely well,  
but that curve-fit handles *new* points very poorly

-> general issue of model complexity

**regularization:** modify error function, to place a “penalty” on large weight values


$$E = \sum (y_i - \text{predicted } y_i)^2 + \alpha \sum w_i^2$$

$\alpha$  = hyperparameter

sometimes called “ridge regression”


**what value of  $\alpha$  to use ?** -> find best  $\alpha$ , to best predict a “hold-back” dataset

# 0<sup>th</sup> Order Polynomial




*this and following slides, adapted from Bishop (2006), Fig.s 1.4 - 1.6*


# 1<sup>st</sup> Order Polynomial



# 3<sup>rd</sup> Order Polynomial



# 9<sup>th</sup> Order Polynomial




-> over-fitting !

# Over-fitting

how to detect over-fitting:


see how well we can predict an independent (“hold-back”) test dataset:




Root-Mean-Square (RMS) Error:  $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$

# Data Set Size:

$M = 9$ , i.e. fit 9<sup>th</sup> order polynomial






With more data, over-fitting becomes less of a problem.




But, gathering more data can sometimes be difficult or expensive ...

# Regularization



penalize large coefficient values:

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



*adapted from Bishop (2006), Fig.s 1.4, 1.7, 1.8*



( DEMO: system identification with gradient descent )

# Optional supplementary reading:

Theoretical Neuroscience Peter Dayan & L.F. Abbott; MIT Press, 2001

Chapter 8: Plasticity and Learning - e.g., 8.4, "Supervised Learning"

Chapter 10: Representational Learning

Pattern Recognition and Machine Learning Christopher Bishop

suppl materials -> <http://research.microsoft.com/en-us/um/people/cmbishop/PRML/index.htm>

## journal articles:

•Wu MCK, David SV, Gallant JL (2006) Complete functional characterization of sensory neurons by system identification. Ann Rev Neurosci 29:477-505.

•Nishimoto S, Gallant JL (2011) A three-dimensional spatiotemporal receptive field model explains responses of area MT neurons to naturalistic movies. J Neurosci 31:14551-14564.

•Talebi V, Baker CL (2012) Natural versus synthetic stimuli for estimating receptive field models: A comparison of predictive robustness. J Neurosci 32:1560-1576.

related courses at McGill: “Machine Learning” (COMP-652) - Doina Precup

## On-line resources:

NetLab - <http://www.ncrg.aston.ac.uk/netlab/>

STRFlab - <http://strflab.berkeley.edu/>

Bruno Olshausen's course (UCBerkeley): [redwood.berkeley.edu/wiki/VS298:\\_Neural\\_Computation](http://redwood.berkeley.edu/wiki/VS298:_Neural_Computation)

Michael Jordan's course (UCBerkeley): [www.cs.berkeley.edu/%7Easimma/294-fall06/](http://www.cs.berkeley.edu/%7Easimma/294-fall06/)

Andrew Moore's slides (CMU): [www.autonlab.org/tutorials/](http://www.autonlab.org/tutorials/)

Wikipedia pages: Neural Networks, Machine Learning, ...